

- [Guided practice on MongoDB Atlas](#)
 - [Many ways to work with MongoDB](#)
 - [Atlas](#)
 - [Important](#)
 - [Then](#)
 - [Querying in MongoDB](#)
 - [Moaaar dataaaa!](#)
 - [Connect via your language](#)
 - [Connect In Python](#)
 - [Use the terminal with `mongosh`](#)
 - [Filtering](#)
 - [Cursor](#)
 - [Projection](#)
 - [Exercises](#)
 - [Number of documents](#)
 - [Your turn](#)
 - [conclusion](#)
 - [going further](#)

Guided practice on MongoDB Atlas

In this document

- we work on MongoDB Atlas a hosted service of MongoDB
- create a database, load some docs and do a few queries
- then move on to working with python or the terminal and a more complex database

At the end of this session you should start to have a good grasp on MongoDB querying language.

Goals:

- You can connect to a MongoDB Atlas server
- You can create a database and insert documents
- You can build somewhat complex queries over a MongoDB database

Many ways to work with MongoDB

On your local machine

- MongoDB Community Edition = Self-hosted, runs on your computer
- MongoDB Compass - the GUI tool for visualization
- `mongosh` the CLI in the terminal

You can also work with your favorite scripting language : python, node.js, go, ruby, PHP, Java, etc ...

Today we work on the hosted version Atlas.

MongoDB Atlas provides Cloud-hosted service with both free and paid tiers.

Atlas gives us a hosted cluster on which we can create a database, import some sample collections and understand how to do CRUD operations in MongoDB. You do not need to install anything.

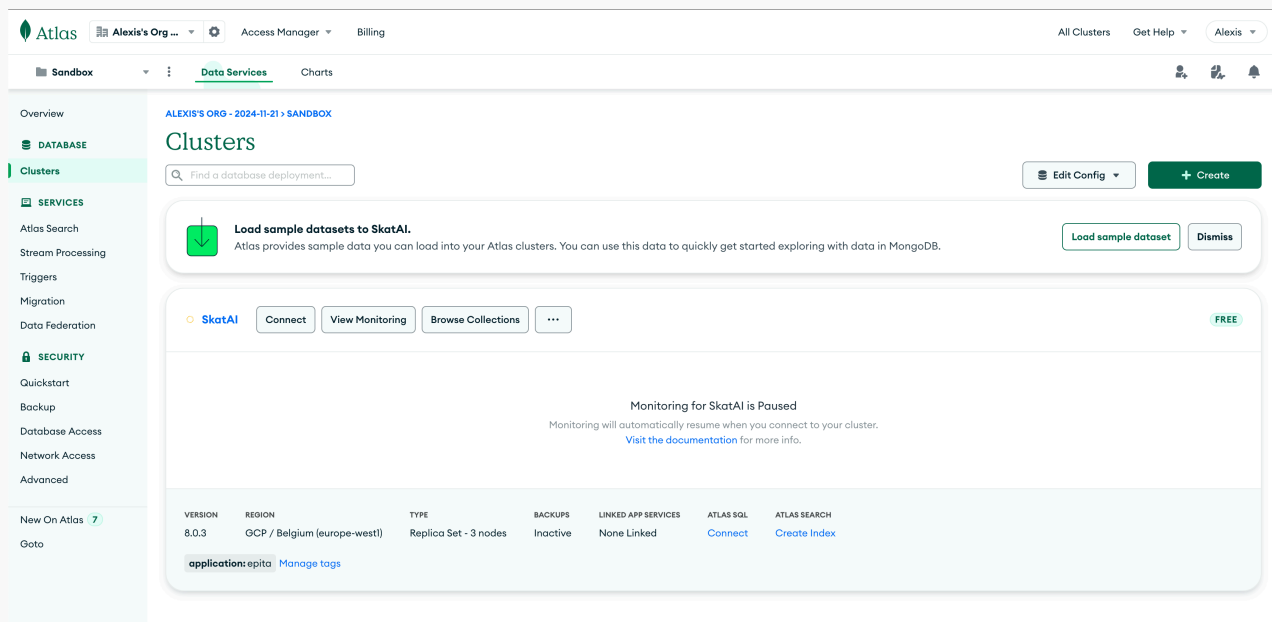
CRUD stands for Create, Read, Update, and Delete : the four basic operations for a DBMS

Atlas

Let's start by creating an account on Atlas <https://www.mongodb.com/cloud/atlas/register>

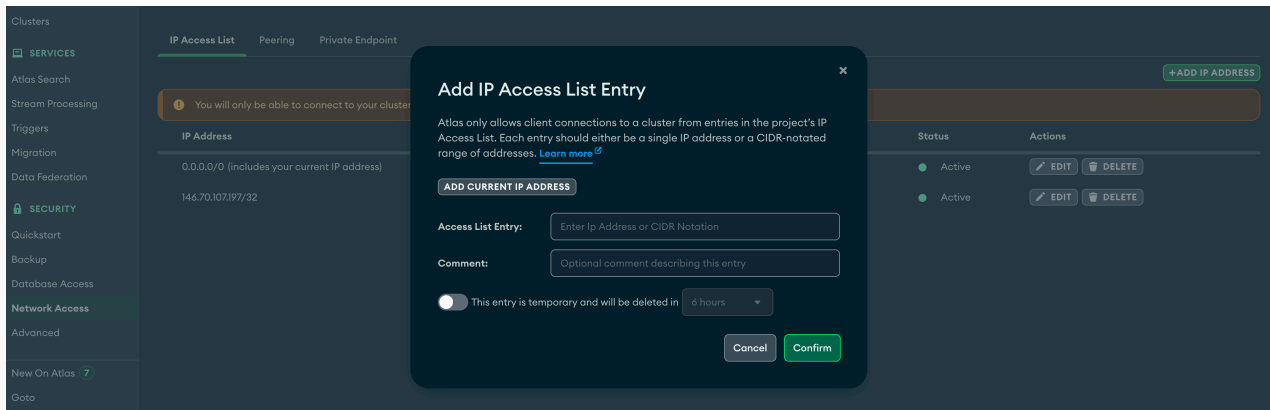
Then:

- Create a cluster and a project (I called my project `Sandbox`)



Important

- **Copy your database password.** It's best to set it as an environment variable or in a `.env` file (for python).
- Check that your **IP** is added to the list, visit `Network Access` (left navigation)
 - allow access from any IP using 0.0.0.0/0 (not recommended for production).



Then

- Go to clusters
- Create a database: add name `songsdb` and collection name `songs`,
- Insert a Document

```
JavaScript
{
  "_id":{
    "$oid":"6745ab6f0e0bbdab062667c7"
  },
  "title": "Happy",
  "artist": "Pharrell Williams",
  "year": 2013,
  "mood": "joyful"
}
```

Insert Document

To collection music

VIEW



```
1  {
2    "_id": {
3      "$oid": "6745ab6f0e0bbdab062667c7"
4    },
5    "title": "Happy",
6    "artist": "Pharrell Williams",
7    "year": 2013,
8    "mood": "joyful"
9  }
10
```



Cancel

Insert

And insert another one

```
"title": "Highway to hell",
"artist": "AC/DC",
"year": 1981,
"mood": "energetic"
```

JavaScript

So we have 2 documents !

Let's explore !!! This is exciting 🤪🤪🤪 !

Querying in MongoDB

In MongoDB querying comes down to writing JSON

json	query
<code>{}</code>	returns all the documents
<code>{ field : value }</code>	where field = value
<code>{ field : { \$lt : value } }</code>	where field <= value

So if we want to find all the documents in our songs collection, simply write '{}' in the query field

This returns our 2 songs.

The screenshot shows the MongoDB Compass interface for the 'musicdb.music' collection. The top navigation bar includes 'Find', 'Indexes', 'Schema Anti-Patterns', 'Aggregation', and 'Search Indexes'. Below the navigation bar, there's a 'Filter' field containing an empty JSON object '{}'. To the right of the filter field are buttons for 'Reset', 'Apply', and 'Options'. Below the filter field, the 'QUERY RESULTS: 1-2 OF 2' are displayed. The first result is a document with the following fields: '_id' (ObjectId), 'title' ('Happy'), 'artist' ('Pharrell Williams'), 'year' (2013), and 'mood' ('joyful'). The second result is a document with the following fields: '_id' (ObjectId), 'title' ('Highway to hell'), 'artist' ('AC/DC'), 'year' (1981), and 'mood' ('energetic').

And the query `{ year : { $lt : 2000 } }` returns the song that was released before 2000.

The screenshot shows the MongoDB Compass interface for the 'musicdb.music' collection. The top navigation bar includes 'Find', 'Indexes', 'Schema Anti-Patterns', 'Aggregation', and 'Search Indexes'. Below the navigation bar, there's a 'Filter' field containing the query '{ year : { \$lt : 2000 } }'. To the right of the filter field are buttons for 'Reset', 'Apply', and 'Options'. Below the filter field, the 'QUERY RESULTS: 1-1 OF 1' are displayed. The single result is a document with the following fields: '_id' (ObjectId), 'title' ('Highway to hell'), 'artist' ('AC/DC'), 'year' (1981), and 'mood' ('energetic').

Mooaar dataaaa!

This is all nice but we need more data to play with.

Let's import the ATLAS sample datasets.

Goto `Clusters > dots > load sample dataset`

Clusters

[Connect](#)
[View Monitoring](#)
[Browse Collections](#)
[Edit Configuration](#)
[Command Line Tools](#)
[Load Sample Dataset](#)
[Terminate](#)


Enhance Your Experience

For production throughput and richer metrics, upgrade to a dedicated cluster now!


[Upgrade](#)

● R 0.07

● W

Last 5 hours

0.2/s



18.0



Then click on your project name (for me Sandbox) and [Browse collections](#) to view the available databases

You get a new database [movies_mflix](#) with 5 collections.

sample_mflix							
LOGICAL DATA SIZE: 115.94MB		STORAGE SIZE: 102.46MB		INDEX SIZE: 19.84MB		TOTAL COLLECTIONS: 6	
Collection Name	Documents	Logical Data Size	Avg Document Size	Storage Size	Indexes	Index Size	Avg Index Size
comments	41079	11.14MB	285B	6.45MB	1	1.91MB	1.91MB
embedded_movies	3483	71.9MB	21.14KB	75.45MB	1	204KB	204KB
movies	21349	32.54MB	1.56KB	20.34MB	2	17.54MB	8.77MB
sessions	1	540B	540B	20KB	2	40KB	20KB
theaters	1564	341.63KB	224B	164KB	2	112KB	56KB
users	185	28.88KB	160B	36KB	2	48KB	24KB

Look at the [embedded_movies](#) collection that has 1525 documents and notice the structure of a document

A document is a JSON record.

It can have:

- **nested JSON** : check out the nested dictionaries [imdb](#) and the [tomatoes](#) fields.
- **arrays**: check out genres, cast, languages, writers, ...

The primary key of a collection is always "_id".

JavaScript

```
{
  "_id": {
    "$oid": "573a1390f29313caabcd5293"
  },
  "plot": "Young Pauline is left a lot of money when her wealthy uncle dies. H",
  "genres": [
    "Action"
  ],
  "runtime": {
    "$numberInt": "199"
  },
  "cast": [
    "Pearl White",
    "Crane Wilbur",
    "Paul Panzer",
    "Edward José"
  ],
  "num_mflix_comments": {
    "$numberInt": "0"
  },
  "poster": "https://m.media-amazon.com/images/M/MV5BMzgxDk1Mzk2Ml5BMl5BanBnX",
  "title": "The Perils of Pauline",
  "fullplot": "Young Pauline is left a lot of money when her wealthy uncle die",
  "languages": [
    "English"
  ],
  "released": {
    "$date": {
      "$numberLong": "-1760227200000"
    }
  },
  "directors": [
    "Louis J. Gasnier",
    "Donald MacKenzie"
  ],
  "writers": [
    "Charles W. Goddard (screenplay)",
    "Basil Dickey (screenplay)",
    "Charles W. Goddard (novel)",
    "George B. Seitz",
    "Bertram Millhauser"
  ],
  "awards": {
    "wins": {
      "$numberInt": "1"
    },
    "nominations": {
      "$numberInt": "0"
    }
  }
}
```

```

    },
    "text": "1 win."
  },
  "lastupdated": "2015-09-12 00:01:18.647000000",
  "year": {
    "$numberInt": "1914"
  },
  "imdb": {
    "rating": {
      "$numberDouble": "7.6"
    },
    "votes": {
      "$numberInt": "744"
    },
    "id": {
      "$numberInt": "4465"
    }
  },
  "countries": [
    "USA"
  ],
  "type": "movie",
  "tomatoes": {
    "viewer": {
      "rating": {
        "$numberDouble": "2.8"
      },
      "numReviews": {
        "$numberInt": "9"
      }
    },
    "production": "Pathé Frères",
    "lastUpdated": {
      "$date": {
        "$numberLong": "1441993579000"
      }
    }
  }
}

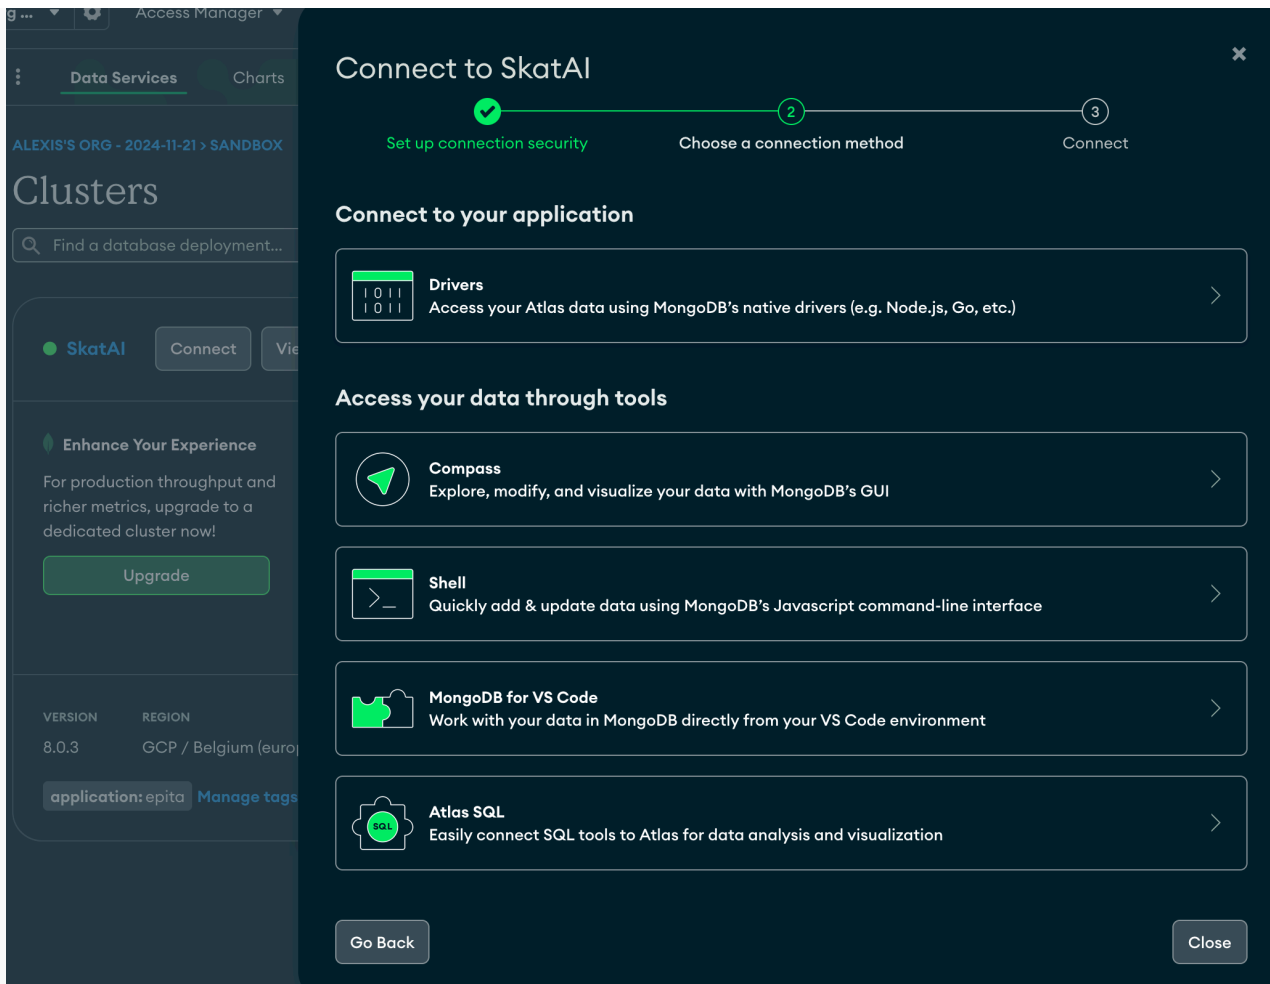
```

The ATLAS UI is great but I'd rather work with code than a UI.

Let's switch to python or `mongosh` to explore that movies database and learn how to query in MongoDB.

Connect via your language

Atlas allows you to connect to your cluster



Connect In Python

We need the [pymongo](#) package

```
pip install pymongo
```

Bash

The connection string is

```
connection_string = "mongodb+srv://alexis:<db_password>@skatai.w932a.mongodb.net"
```

Python

Note: it's best to put the connection string as an environment variable (`$MONGO_ATLAS_URI`), for instance in a `.env` file. then load it with

```
import os
from dotenv import load_dotenv

load_dotenv()

MONGO_ATLAS_URI = os.getenv('MONGO_ATLAS_URI')
```

Python

And then

```
from pymongo import MongoClient

client = MongoClient(MONGO_ATLAS_URI)
```

Python

Once we have a client we can connect to the database

```
db = client["sample_mflix"]
```

Python

and instantiate a collection object

```
collection = db["movies"]
```

Python

The collection object is of the class `pymongo.synchronous.collection.Collection` and has many methods:

<code>aggregate()</code>	<code>count_documents()</code>	<code>create_search_indexes()</code>	<code>distinct()</code>	<code>drop_search_index()</code>	<code>find_one_and_delete()</code>
<code>aggregate_raw_batches()</code>	<code>create_index()</code>	<code>database</code>	<code>drop()</code>	<code>estimated_document_count()</code>	<code>find_one_and_replace()</code>
<code>bulk_write()</code>	<code>create_indexes()</code>	<code>delete_many()</code>	<code>drop_index()</code>	<code>find()</code>	<code>find_one_and_update()</code>
<code>codec_options</code>	<code>create_search_index()</code>	<code>delete_one()</code>	<code>drop_indexes()</code>	<code>find_one()</code>	<code>find_raw_batches()</code>
<code>function(pipeline: 'Pipeline', session: 'Optional[ClientSession]'=None, let: 'Optional[Mapping[str, Any]]'=None, comment: 'Optional[Any]'=None, **kwargs: 'Any')</code>					

Use the terminal with `mongosh`

How to install `mongosh`: <https://www.mongodb.com/docs/mongodb-shell/install/>

python: the quotes

```
db.movies.find(
    {"runtime": {"$gt" : 180}}, // Filter on movie duration
    { "_id": 0, "title": 1, "runtime": 1, "imdb.rating": 1 } // Projection to include title and runtime and rating
)
```

Python

`mongosh`: no quotes

```
db.movies.find(
  {runtime: {$gt : 180}}, // Filter on movie duration
  { _id: 0, title: 1, runtime: 1, "imdb.rating": 1 } // Projection to include title and runtime and rating
)
```

JavaScript

Filtering

<https://www.mongodb.com/docs/manual/reference/glossary/>

The JSON that specifies the filtering arguments is called a **query predicate**. It is an expression that returns a boolean indicating whether a document matches the specified query.

For example, `{ title: { $eq: "Dil Se" } }`, which returns documents that have a field "title" whose value is the string "Dil Se".

An empty query predicate (`{ }`), the query returns all documents in the collection.

Main functions on a collections

function	returns
<code>find()</code>	all the documents
<code>find_one()</code>	the 1st document
<code>distinct("<field>")</code>	list of distinct values for the <code><field></code>
<code>count_documents({})</code>	number of documents for the collection or returned by the filter in the query predicate

Note also

- `find_one_and_replace()`
- `find_one_and_update()` and
- `delete_many()`
- `delete_one()`
- `drop_index()`

Note: You can also just query the collection directly from the client with `db.<collection_name>.find()`

```
collection.find({})  
# or  
db.movies.find({})
```

Bash

You can chain these methods with limit and sort

```
db.movies.find( {runtime: {$gt: 120}} ).limit(3)
```

Bash

Cursor

The returned result is a **cursor**.

```
cursor = db.movies.find({})
```

Python

A cursor is a pointer over a MongoDB query set of results.

Projection

In database speak, **projecting** means selecting a subset of all the fields.

In SQL, you simply list the column names

```
select genres, plot from movies;
```

SQL

In MongoDB, specify in a json object which field you want to see right after the query predicate

```
db.movies.find(
  {runtime: {$gt : 180}}, // Filter on movie duration
  { _id: 0, title: 1, runtime: 1, "imdb.rating": 1 } // Projection to include
)
```

JavaScript

Projection: ({ _id: 0, title: 1, runtime: 1, "imdb.rating": 1 })

- `title: 1`: Includes the title field.
- `runtime: 1`, includes the runtime
- `"imdb.rating": 1`: Includes the `imdb.rating` field.
- `_id: 0`: Excludes the `_id` field from the result (default is 1 if not specified).

The query returns

```
{ runtime: 240, title: 'Napoleon', imdb: { rating: 7.4 } },
{ runtime: 281, title: 'Les Misérables', imdb: { rating: 7.9 } },
{ runtime: 245, title: 'Flash Gordon', imdb: { rating: 7.3 } },
{ runtime: 238, title: 'Gone with the Wind', imdb: { rating: 8.2 } },
```

JavaScript

Exercises

Let's run a few queries in python on the movies database

```
import os
from pymongo import MongoClient

connection_string = os.getenv('MONGO_ATLAS_URI')
client = MongoClient(connection_string)
db = client["sample_mflix"]
```

Python

Then

- Retrieve all movies that have the genre "Action". Only get the title and genres.

Python

```
cursor = db.movies.find(
    {"genres": "Action"}, # Filter: movies with 'Action' in the genres array
    {"_id": 0, "title": 1, "genres": 1} # Projection: include title and genres,
)
```

to see the results

Python

```
for movie in cursor:
    print(movie)
```

with `mongosh`:

Python

```
cursor = db.movies.find(
    {"genres": "Action"}, # Filter: movies with 'Action' in the genres array
    {"_id": 0, "title": 1, "genres": 1} # Projection: include title and genres,
)
```

Number of documents

Fastest way is to use `count_documents`

Python

```
count = db.movies.count_documents({"imdb.rating": {"$gt": 8.0}})
```

Note: in python you can clone the cursor to get it's length and the number of returned documents. Cloning the cursor does not consume it

Python

```
len(list(cursor.clone()))
```

Your turn

with

Python

```
cursor = db.movies.find( filter, projection).limit(5)
for movie in cursor:
    print(movie)
```

Write the filter and projection for the following queries and return also the number of documents with `db.movies.count_documents(filter)`

- use projection to only return relevant fields or at minimum "title"
- limit the results to 5 documents

1. Find Movies with an IMDb Rating Greater Than 8
 - filter: `{"imdb.rating": {"$gt": 8}}`
 - projection: `{"_id": 0, "title": 1, "imdb.rating": 1}`
2. Movies Released After 2000
3. Movies with Specific Directors: "Christopher Nolan". Show the title, director and year
4. Retrieve movies with a `tomatoes.viewer.rating > 4.0`, showing the title and viewer rating.
5. Find movies that contain "Comedy" and "Drama" in the `genres` array. use `{$all: [list of genres]}`
6. Combine Query with Sorting: Retrieve the top 5 movies with the highest IMDb rating, showing title and rating. (you need to only retrieve `imdb.rating` with data type `double`)
7. Query Movies with a Range of Years: Retrieve movies released between 1990 and 2000, showing the title and year.
8. Query Movies with Missing Fields: Find movies where the `fullplot` field does not exist. use `$exists`.
9. find all the distinct genres
 - use `db.movies.distinct("genres")`
10. movies with at least 2 genres
 - use `{"genres": {$size: 2}}`
11. Movies with genre Action, after 1950 with imdb ratings > 8, sort by year desc, imdb rating desc
 - use `{"year": {"$gt": 1950}, "imdb.rating": {"$gt": 8}, "genres": "Action" }`
12. Movies with both genres : Action and Drama
 - `$and: [{"genres": "Action"}, {"genres": "Drama"}]`
13. Movies with either Action or Drama
 - `$or: [{"genres": "Action"}, {"genres": "Drama"}]`
14. Movies after 1950 with either `imdb.rating > 0` or `awards.wins > 5`
 - use: `{"year": {"$gt": 1950}, $or: [{"imdb.rating": {"$gt": 8}}, {"awards.wins": {"$gt": 5}}] }`

etc etc

conclusion

In this session, you have learned:

- DBMS & history of databases
- An overview of the different types of databases
- relational vs non relation database
- flexible schema in NoSQL databases

- Why and when to choose MongoDB over SQL

And you practiced:

- Setting up MongoDB Atlas, a cloud-hosted database service, including cluster creation and security configuration
- Writing basic MongoDB queries using JSON format:
 - Basic syntax: `{}` for all documents, `{field: value}` for equality, `{field: {$lt: value}}` for comparisons
 - How to query nested fields and arrays in complex documents
- Connecting to MongoDB Atlas using Python and `pymongo`:
 - Using basic operations: `find()`, `find_one()`, `distinct()`, `count_documents()`
 - Implementing projections to select specific fields
- Working with sample datasets (particularly movies database) to practice:
 - Filtering and sorting data
 - Working with nested fields
 - Using operators like `$gt`, `$lt`, `$all`, `$exists`
 - Writing combined queries with multiple conditions

In the next session, we deep dive into MongoDB and look at more complex ways to query the data using **aggregation pipelines**. We also tackle schema design and validation.

going further

For next time, you can

- explore more of the Atlas sample databases and practice your querying skills
- follow <https://www.mongodb.com/docs/languages/python/pymongo-driver/current/read/> to get more practice
- There are many free courses and tutorials in the mongoDB university
 - [Intro to MongoDB](#)
 - [CRUD in python](#)

and many more